

dbmv – Database Schema Conversion Tool

1.0 Dbmv.py

- 1.1 Source Schemas
- 1.2 Target Schemas
- 1.3 Syntax
- 1.4 Supported Python Drivers
- 1.5 Logging
- 1.6 Examples

2.0 Your Support Options

1.0 Dbmv.py

The dbmv.py tool helps you to move one database schema to another database if you have a "live" connection available to both source and target databases. ODBC is used both to manage metadata about the schemas and to transfer data (if required), so Python-accessible ODBC drivers are required on the execution machine.

1.1 Source Schemas

The source schema can be from any of the following database products:

- ORACLE Mysql
- MICROSOFT SqlServer
- ORACLE
- POSTGRES
- SAP Sybase IQ (V15+)
- TERADATA

Schema description can easily be extended to any database in the destination list. You must add missing description for "tbDefinition", etc., in file "dbmv.xml".

SQL output, field names, and SELECT descriptions must respect the pattern used for other existing databases.

1.2 Target Schemas

The target schema can be any of the following:

- ACTION VectorH, Vector

- ACTIAN Ingres
- EMC Greenplum
- IBM Netezza (UDB, not sure about iseries and Zos)
- IBM Netezza (4.xx ?)
- ORACLE MySQLI
- MICROSOFT SQL Server
- ORACLE
- POSTGRES
- PROGRESS Database
- SAP Sybase IQ (V15+)
- SAP Sybase ASA (Anywhere)
- SAP Sybase ASE
- SAP Hana
- SAP MaxDB
- TERADATA

1.3 Syntax

Without parameters, dbmv.py displays simple help about its syntax.

```
$ ./dbmv.py
```

```
usage: dbmv_new.py [-h] [--src SRC] [--dest DEST] [--loaddata] [--cretab]
                  [--creview] [--creall] [--loaddl] [--loadtest]
                  [--batchsize BATCHSIZE] [--maxrows MAXROWS] [--truncate]
                  [--parfile PARFILE] [--fdelim FDELIM] [--unload]
                  [--translation TRANSLATION] [--quote QUOTE]
                  [--cmdsep CMDSEP] [--charmax CHARMAX] [--creindex]
                  [--ownsrc OWNSRC] [--owntgt OWNTGT] [--add_drop]
                  [--on_error ON_ERROR] [--source_schema SOURCE_SCHEMA]
                  [--target_schema TARGET_SCHEMA] [--unsupported]
                  [--exclude EXCLUDE] [--include INCLUDE] [--tables TABLES]
                  [--insertmode INSERTMODE] [--trial]
                  [--loadmethod LOADMETHOD] [--threads THREADS]
```

Optional arguments are as follows:

-h, --help

Shows the help message.

--on_error=ON_ERROR

Specifies action to take when an error occurs: **continue** or **abort**.

--src=SRC

source DB used as a data source for migration, copy, etc.. For more information, see [Specifying Source and Destination Databases](#).

--dest=DEST

Destination DB used as a destination for migration. Copy , etc. For more information, see [Specifying Source and Destination Databases](#).

--crtab

Creates only tables in the @dest from @source schema.

Note: Use crtab and creall to build a full schema.

--creview

Creates only views in the @dest from @source schema.

Note: Several iterations may be required if nested views exist. Depends on how the source database returns the definitions.

--creindex

Creates only indexes in @dest schema.

--creall

Creates all objects of type INDEX, PRIMARY KEY, REFERENTIAL CONSTRAINTS, and VIEWS in a @dest from @source schema.

Note: creall does **not** create objects of type TABLE; use crtab and creall to build a full schema. Such an operation may perform slower when loading data into the tables in addition to creating all the objects.

--loadddl

Loads DDL from @source to @dest. If not specified, a file containing all objects is created in the current directory. In this case, you can reload.

--add_drop

Drops table at the beginning.

--truncate

Removes existing rows from @dest table.

Default: False

--unsupported

Whether to skip unsupported types or not.

Default: False

--loaddata

Indicates whether to load the table data from @source to @dest. For more information, see [Unload and Load Capabilities](#).

--loadtest

Performs predefined INSERT queries to @dest (used for performance testing).

Default: False

--batchsize=BATCHSIZE

Sets custom batch size for INSERT queries.

Default: 500

Limits: 1 to 10000

--maxrows=MAXROWS

Sets total limit for INSERT queries.

Default: 100000

Limits: 1 to 100000

--parfile=PARFILE

Specifies a file with dbmv parameters. Parameters are stored as a parfile. Parfiles examples can be found in the 'wrk' directory.

--fdelim=FDELIM

Sets field delimiter for data or statements written to file. Used for data unloads.

Default: \t

--unload

Unloads table structure from @source to file as CREATE statements. For more information, see [Unload and Load Capabilities](#).

--translation=TRANSLATION

Uses translation table/constraints/indexes rules in @dest. Use `schema-name` translation pairs for tables, constraints, and indexes.

Example:

```
'''--translation=sname:public,vw_user,postgres,vw_user;  
iscname:public,vw_user,postgres,vw_user; rscname:public,vw_user,postgres,vw_user'''
```

Explanation:

sname: Tables owned by "public" are moved to schema "vw_user", and from "postgres" to "vw_user".

iscname: Indexes owned by "public" are moved to schema "vw_user" and from "postgres" to "vw_user".

rscname: Constraints owned by "public" are moved to schema "vw_user" and from "postgres" to "vw_user".

--quote=QUOTE

Uses specified symbol as quotation symbol for DDL. Example:

```
--quote=''' => CREATE TABLE "my_table"
```

--cmdsep=CMDSEP

Specifies command separator used in DDL file.

Default=';'

Examples:

```
--cmdsep='\g'
```

```
--cmdsep='go'
```

--charmax=CHARMAX

Specifies maximum column precision that is acceptable for translation.

Default: 6400

Limits: 1 to 30000

--ownsrc=OWNSRC

Specifies single source owner.

--owntgt=OWNTGT

Specifies single new target owner.

--source_schema=SOURCE_SCHEMA

Specifies @source schema to process.

--target_schema=TARGET_SCHEMA

Converts @source_schema to the specified schema.

--exclude=EXCLUDE

Excludes specified tables or columns from processing.

--include=INCLUDE

Includes only specific tables or columns for processing.

--tables=TABLENAME

Process only specified tables

--insertmode=INSERTMODE

Specifies mode for inserting data:

bulk - Insert data in chunks (buffered insertion) (Default).

row - Insert data row-by-row, ignoring batchsize.

--trial

Tests the operation (trial mode).

--loadmethod=LOADMETHOD

Specifies number of threads to use when loading data:

serial - 1 thread (Default)

parallel - 4 threads

multitable - CPU dependent

--threads=THREADS

Specifies custom number of threads to use when loading data.

Default: 1

Limits: 1 to 32 (CPU dependent)

1.3.1 Specifying Source and Destination Databases

The --src and --dest arguments are specified as follows:

--src=URL

--dest=URL

where

URL: db_driver[-odbc]://db_host[:port]/dbname?db_login&login_password

And where db_driver can be one of the following driver names:

<u>Driver Name</u>	<u>Product</u>
vector	ACTIAN Vector
vectorh	ACTIAN VectorH
ingres	ACTIAN Ingres
greenplum	EMC Greenplum
netezza	IBM Netezza
mysql	ORACLE MySQL
mssql	MICROSOFT SQL Server
oracle	ORACLE
postgres	POSTGRES
progress	PROGRESS Database
iq	SAP Sybase IQ (V15+)
asa	SAP Sybase ASA
ase	SAP Sybase ASE
hana	SAP Hana
maxdb	SAP MaxDB
teradata	TERADATA

1.3.1.1 About ODBC

Linux:

The path and configuration for the ODBC database driver must be added in file "../etc/driverTools.xml".
(See in-file examples.)

Windows:

ODBC drivers and Datasource must exist. No additional configuration is required. The DSN is provided in place of db_host:

URL : db_driver[-odbc]://DSN[:port]/dbname?db_login&login_password

Note: port and dbname are required but ignored for Windows DSN use case.

When using Azure and Microsoft SQL Server:

--src=mssql-
odbc://your_azure_access.database.windows.net/dbname?db_login&login_password

Note: The .database.windows.net is a generic value provided by Azure.

On all operating systems, Microsoft SQL Server ODBC driver Version 13 or higher is required.

1.3.2 Unload and Load Capabilities

The --unload and --loaddata parameters provide basic facilities for unloading and loading data. These operations are inefficient for massive data but can be useful for small tables (less than 100,000 records).

1.4 Supported Python Drivers

Supported Python drivers are as follows:

<u>Module</u>	<u>Product</u>
pyodbc	ODBC (Used for many databases including Actian Vector, VectorH, Ingres, Microsoft SQL Server on premises, and Azure.)
cx_Oracle	Oracle
Sybase	Sybase ASE
pymssql	MS-SQL
MySQLdb	My SQL
psycopg2, psycopg2.extensions	Postgres
import DB2	DB2

1.5 Logging

The dbmv tool provides JSON logging and is configured to use a relative path by default from the DBMV_HOME/bin folder.

To override this, use the LOG_CFG environment variable. Define the LOG_CFG variable to be the full path to the JSON logging configuration file:

Windows:

```
set LOG_CFG=C:\Actian\dbmv\etc\logging.json
```

Linux:

```
export LOG_CFG=/opt/Actian/dbmv/etc/logging.json
```

1.6 Examples

To run correctly, all parameters must be on one line. For clarity, the parameters are shown here on separate lines.

1. Load all tables for dbo schema, excluding columns not required, specifying the source and destination in the parameter file ms_vw.par.

Note: The image columns and binary datatype columns are automatically ignored for Vector and VectorH destinations.

```
C:\DBMV_HOME> python.exe bin\dbmv_new.py
--quote="
--parfile=C:\DBMV_HOME\wrk\ms_vw.par
--cretab
```

```
--source_schema=dbo
--maxrows=3000
--target_schema=vw_user
--batchsize=700
--loadddl
--add_drop
--loaddata
--exclude="Employees"."Photo","Categories"."Picture"
--threads=10
--maxrows=2000
--creall
```

2. Add additional views if there were load failures.

Note: Load failures can happen when some views rely on others and the migration order is not aligned to the creation of the views. This is a small limitation of dbmv, depending on the source database.

```
C:\DBMV_HOME> python.exe bin\dbmv_new.py
--quote=\"
--parfile=C:\DBMV_HOME\dbmv\wrk\ms_vw.par
--source_schema=dbo
--target_schema=vw_user
--loadddl --exclude="Employees"."Photo","Categories"."Picture"
--creview
```

EXAMPLE of Oracle odbc based source

```
--src=oracle-odbc://SD02/SD02?db_login&login_password
```

EXAMPLE of Vector Based destination

```
--dest=vector-odbc://topcoder64:II/topcoder?db_login&login_password
```

2.0 Your Support Options

Enterprise customers with active maintenance and support contracts have full access to Actian Support, including online use of our case management system and knowledge base, at the Customer Portal (<https://communities.actian.com>).

If you have an active support contract and want to register for access to the Customer Portal, use the enrollment form at <http://supportservices.actian.com/user/register.php>. (Your Account Number ID is required.)

If you do not have a support agreement for Actian Corporation products and are interested in purchasing support, contact us at sales@actian.com.

For more information about support options, visit <https://www.actian.com/support-services/>.